




You say big data
I say fast data

Nilanjan Raychaudhuri

@nraychaudhuri






Big data

Big Data

Word for storing, managing and making money of very large dataset



fast data

<http://pichost.me/1628868/>

Fast Data

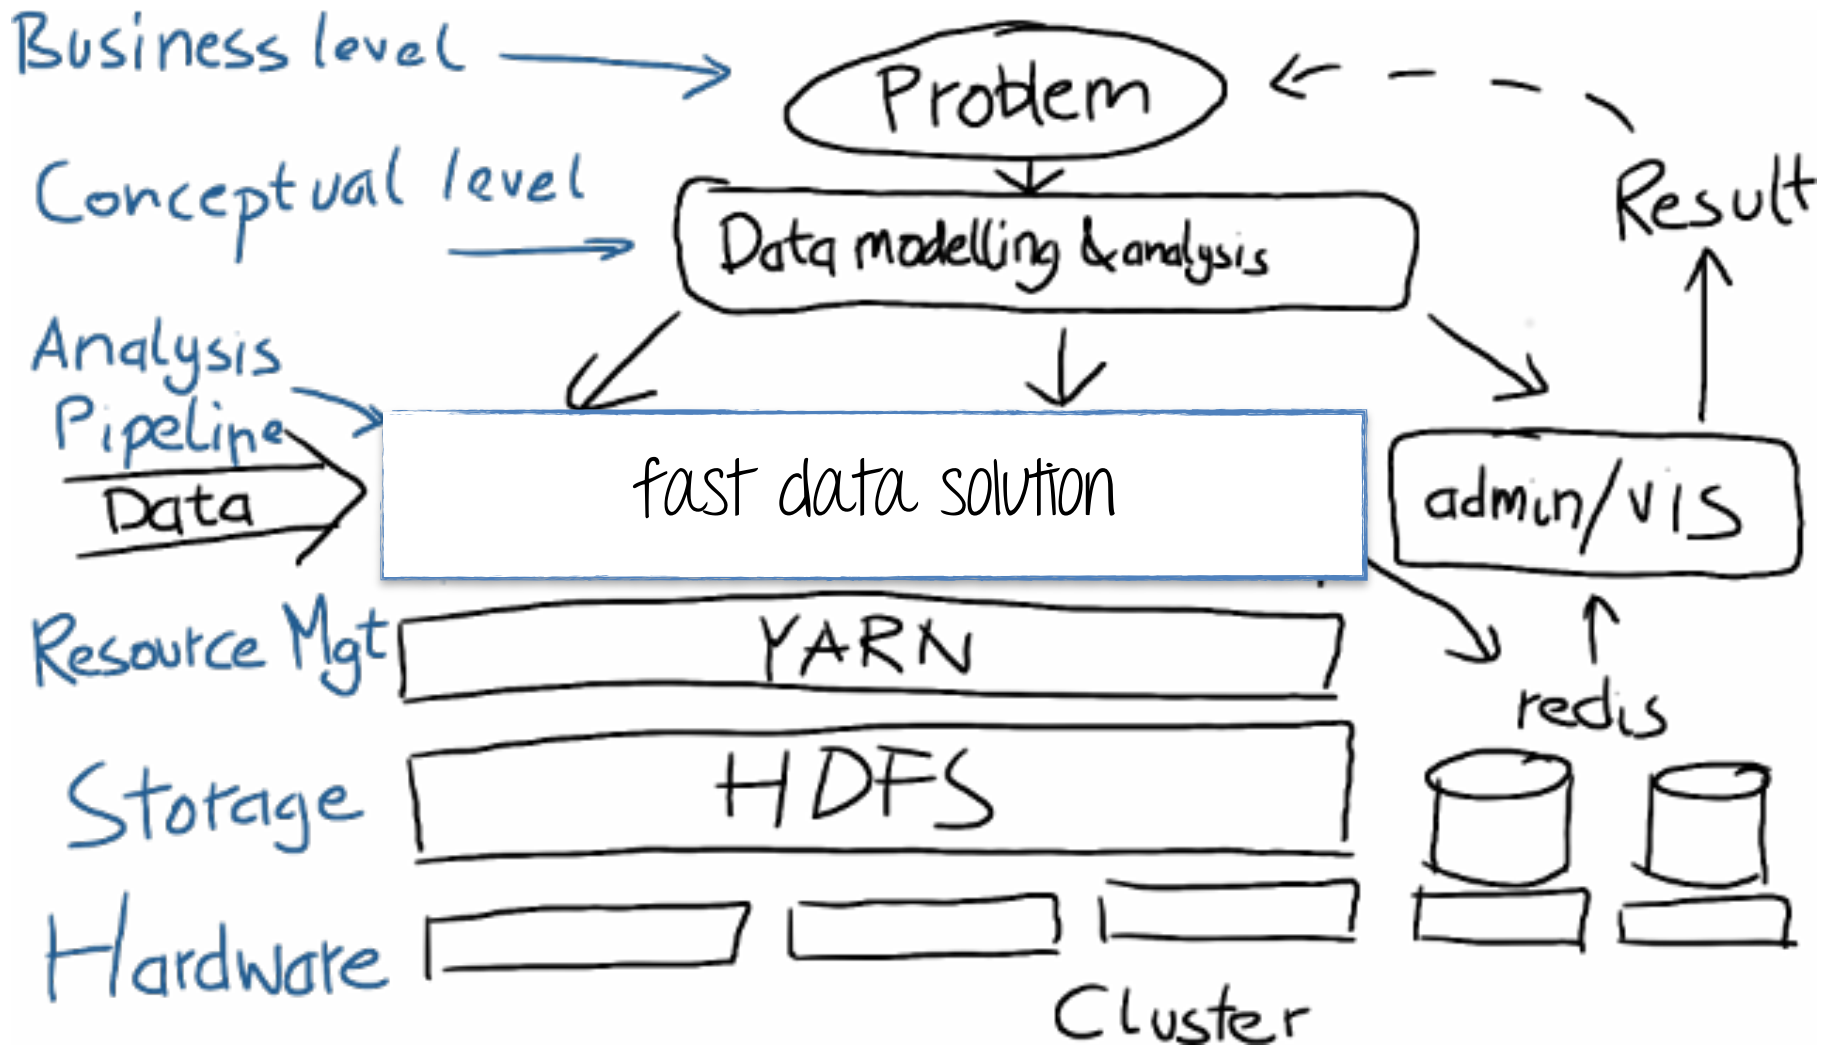
Word for using fast, real time analytics, online machine learning for profit

Definition

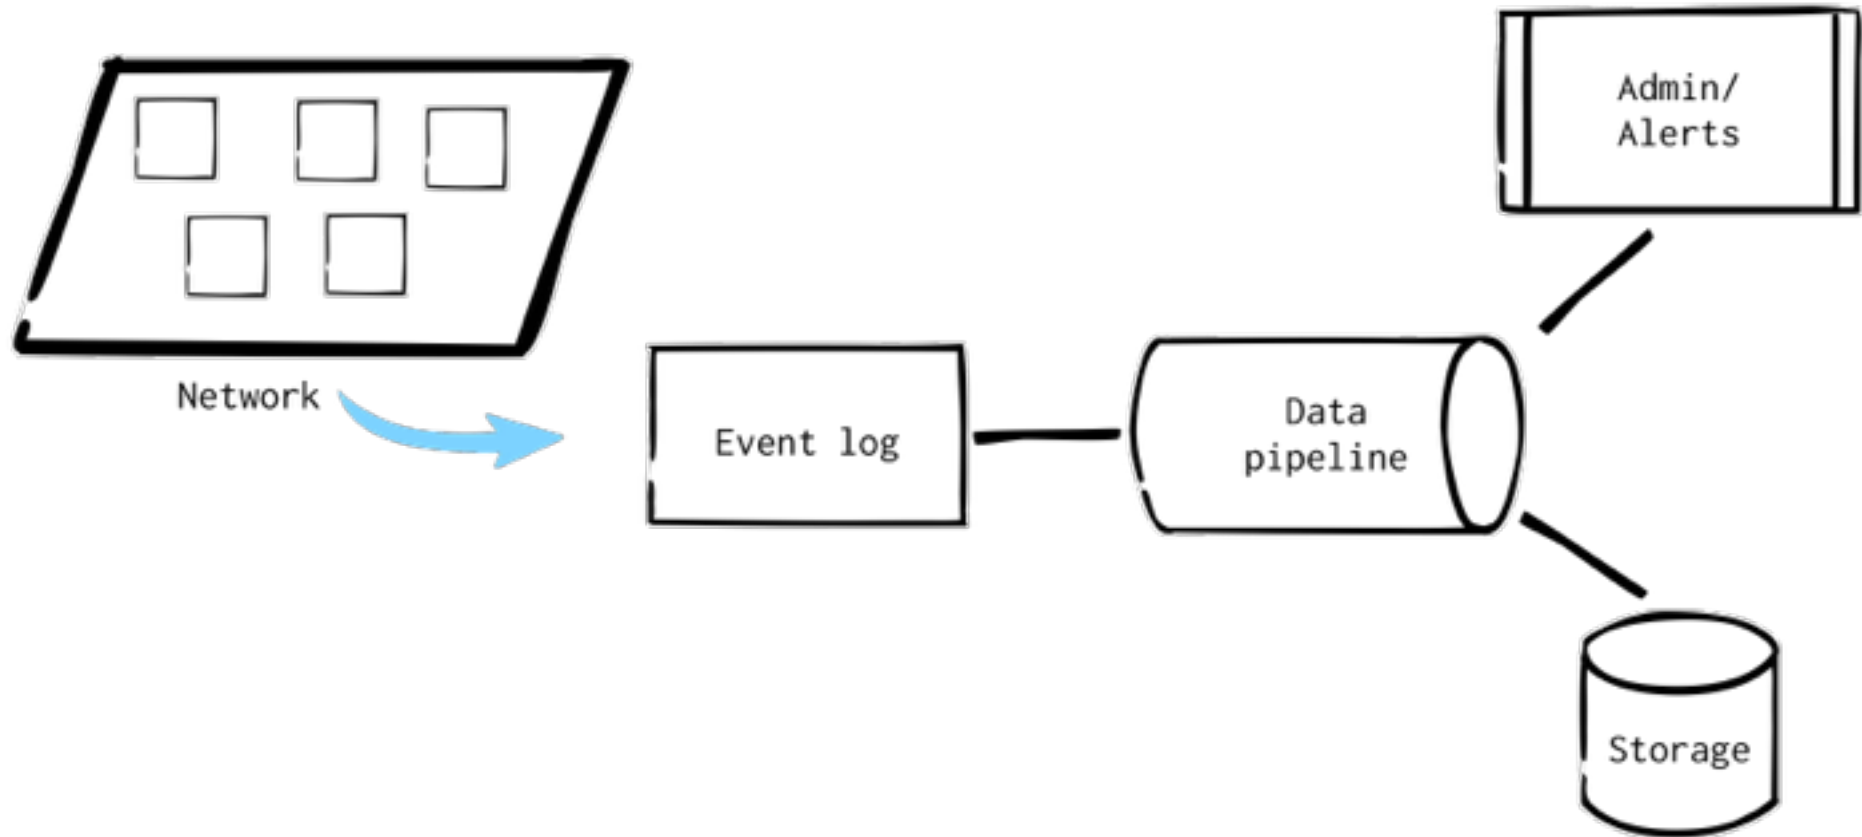
The phrase *Fast Data* captures this range of new systems and approaches, which balance various tradeoffs to deliver timely, cost-efficient streaming data processing.

"You can't expect the value of data to just appear out of thin air. Data isn't fissile material. It doesn't spontaneously reach critical mass and start producing insights."

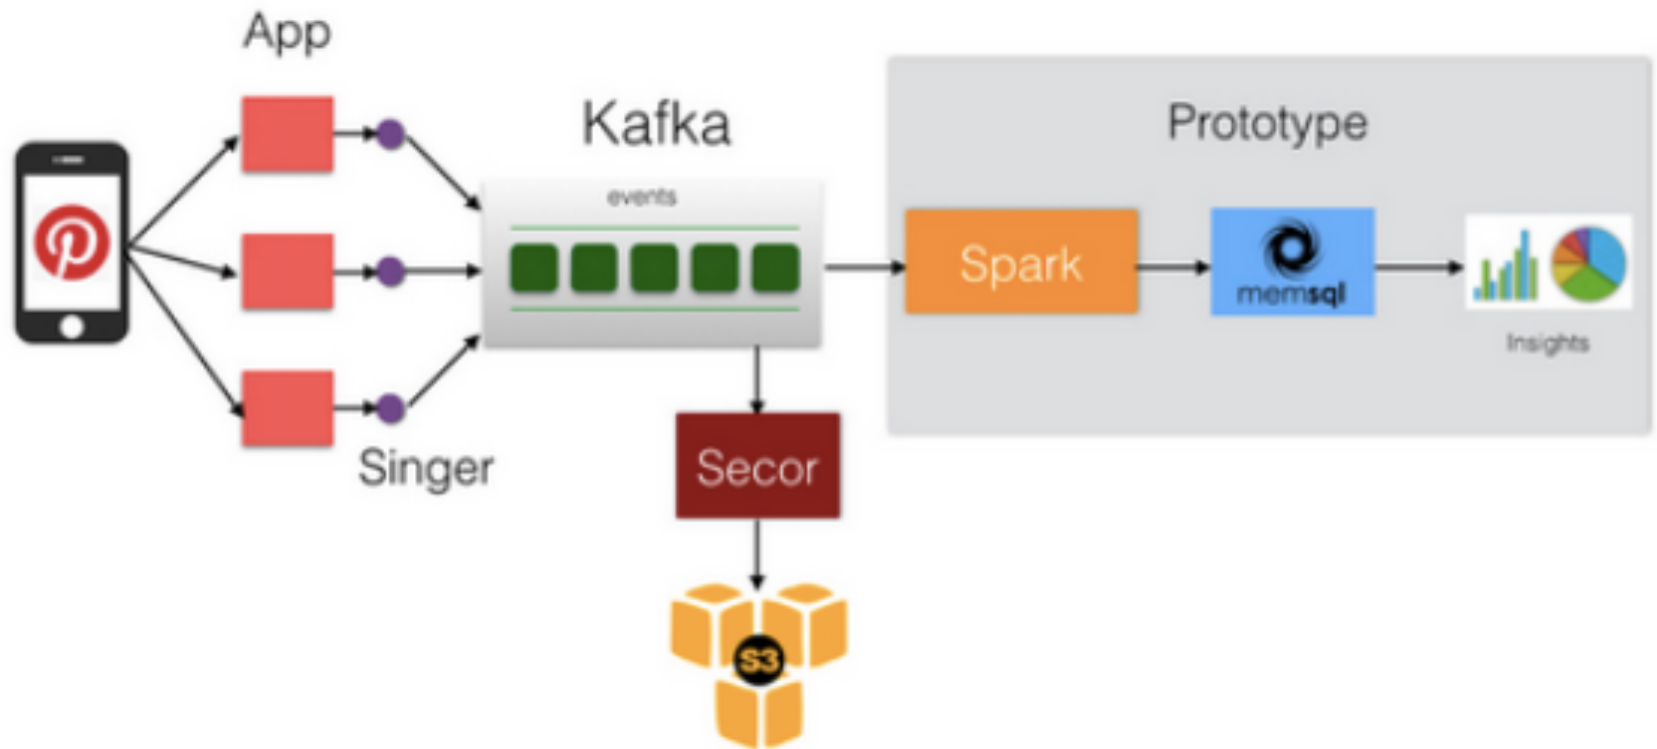
- Marko Karppinen



Detecting network intrusion

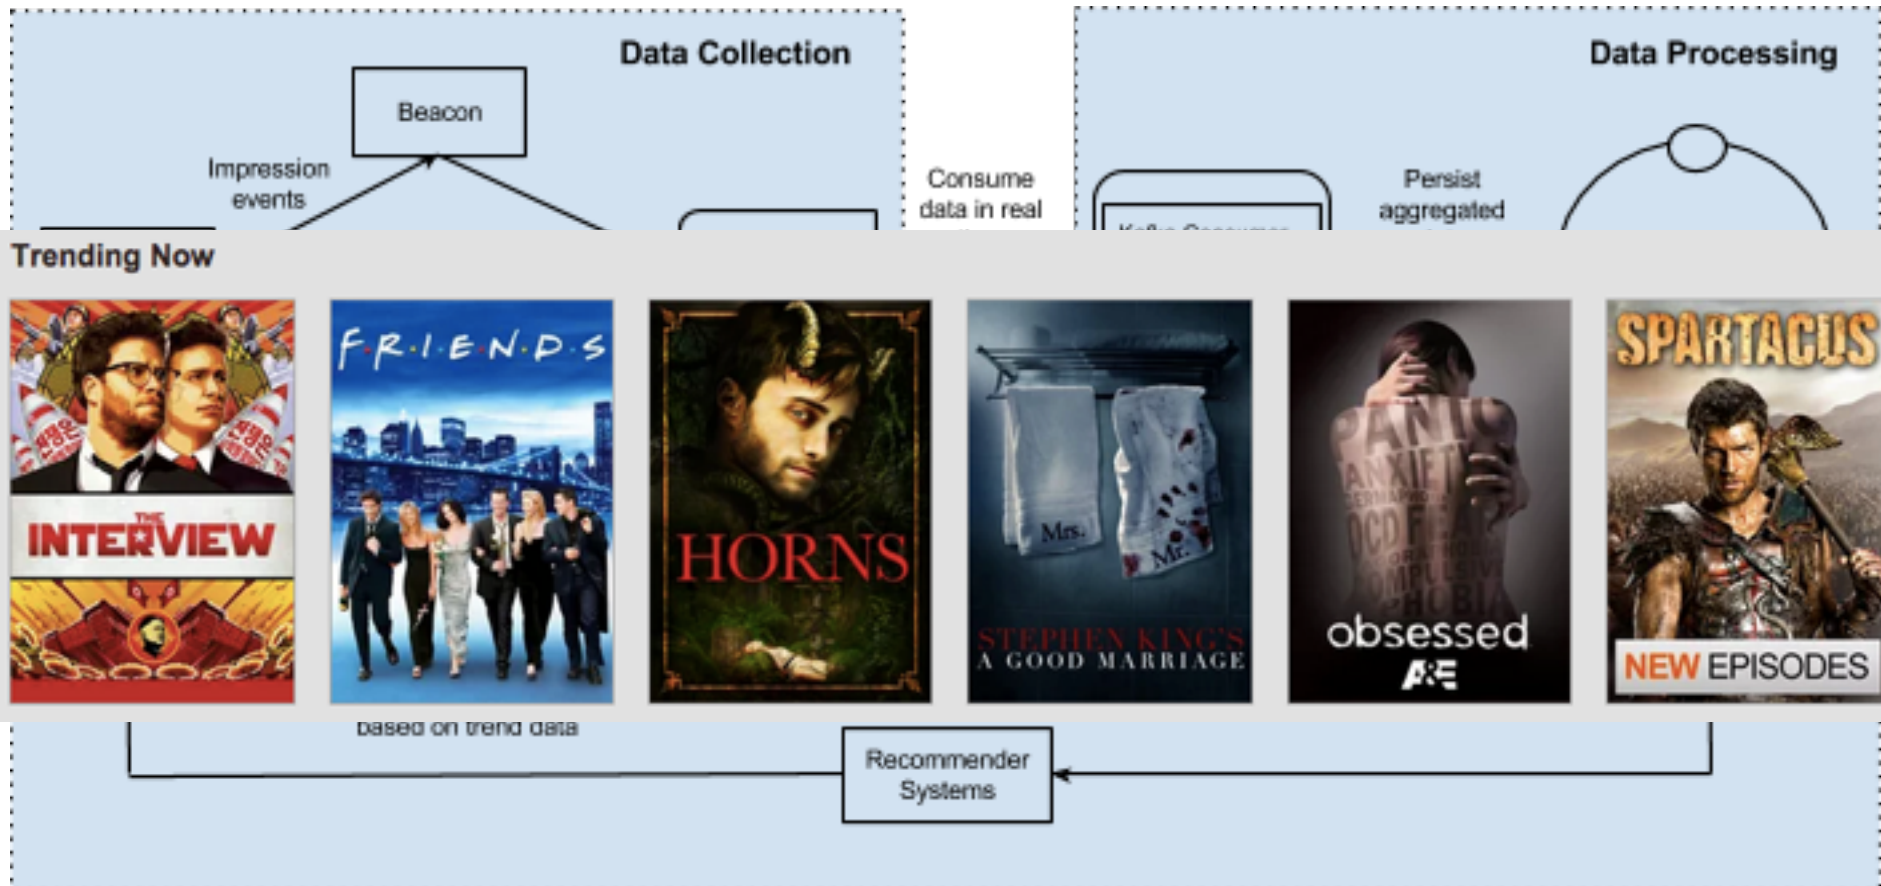


Faster data analysis @ Pinterest



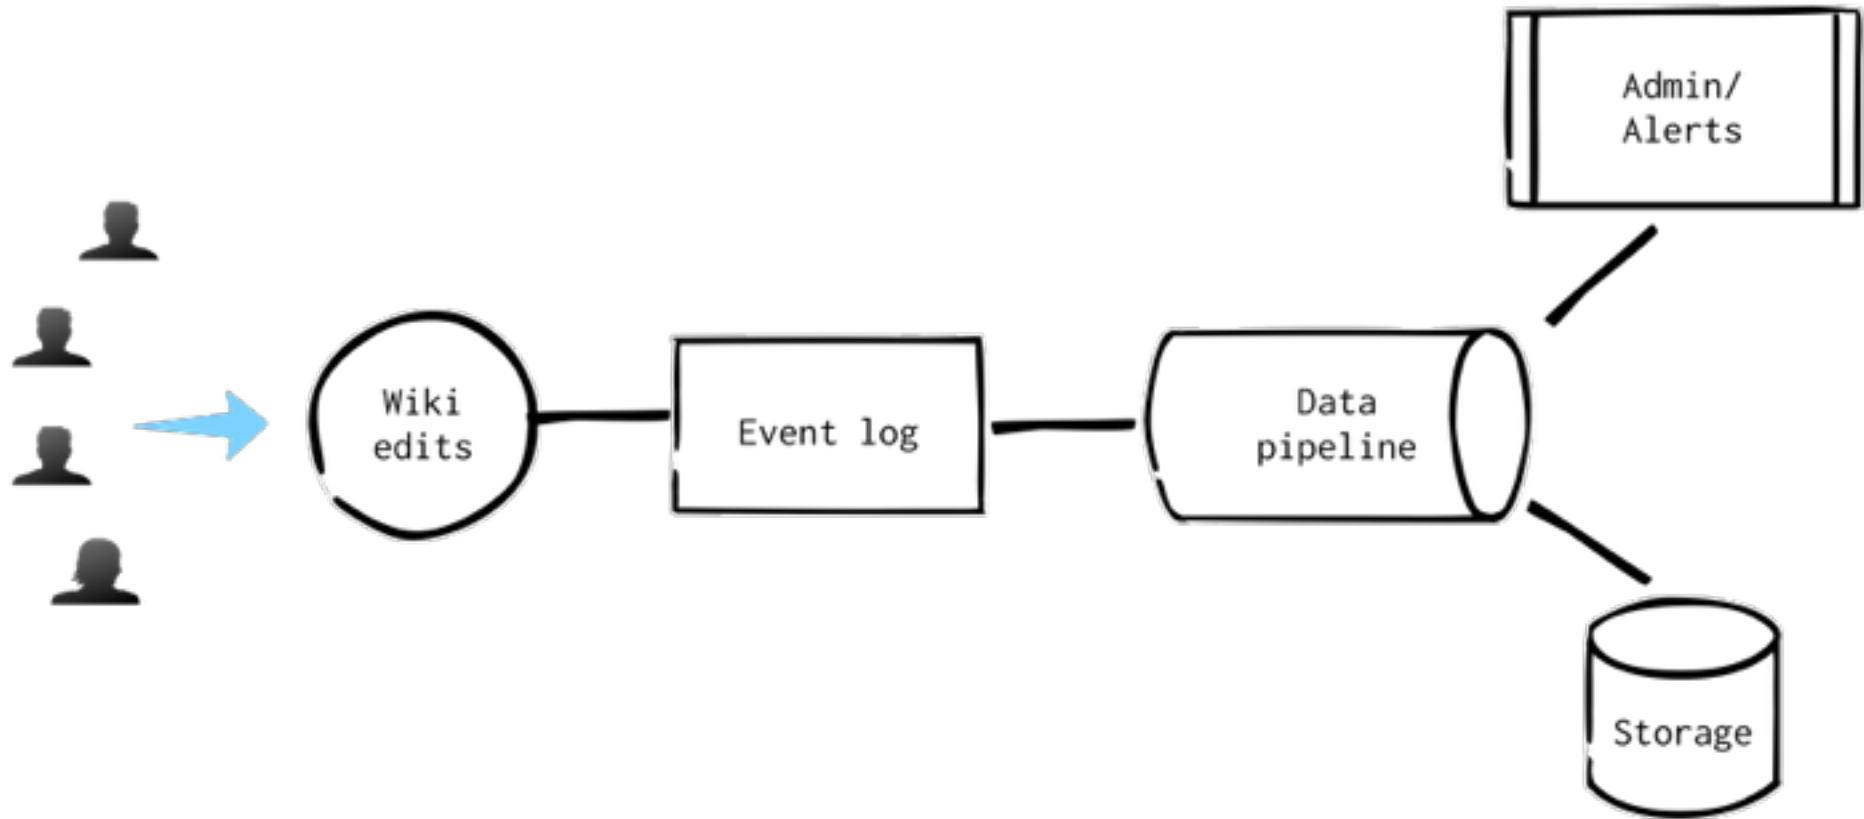
<http://engineering.pinterest.com/post/111380432054/real-time-analytics-at-pinterest>

Recommendation engine @ Netflix



<http://techblog.netflix.com/2015/02/whats-trending-on-netflix.html>

Predict breaking news




What big data/fast data has to do with
functional programming?

Functional programming is the killer Paradigm for Data Apps

- Dean Wampler, Ph.D

Dataflow programming



```
sparkContext.textFile("/path/to/input")
  .map { line =>
    val array = line.split(",", 2)
    (array(0), array(1))
  }.flatMap {
    case (id, contents) => toWords(contents).map(w => ((w, id), 1))
  }.reduceByKey {
    (count1, count2) => count1 + count2
  }.map {
    case ((word, path), n) => (word, (path, n))
  }.groupByKey
  .map {
    case (word, list) => (word, sortByCount(list))
  }.saveAsTextFile("/path/to/output")
```

What is streaming?

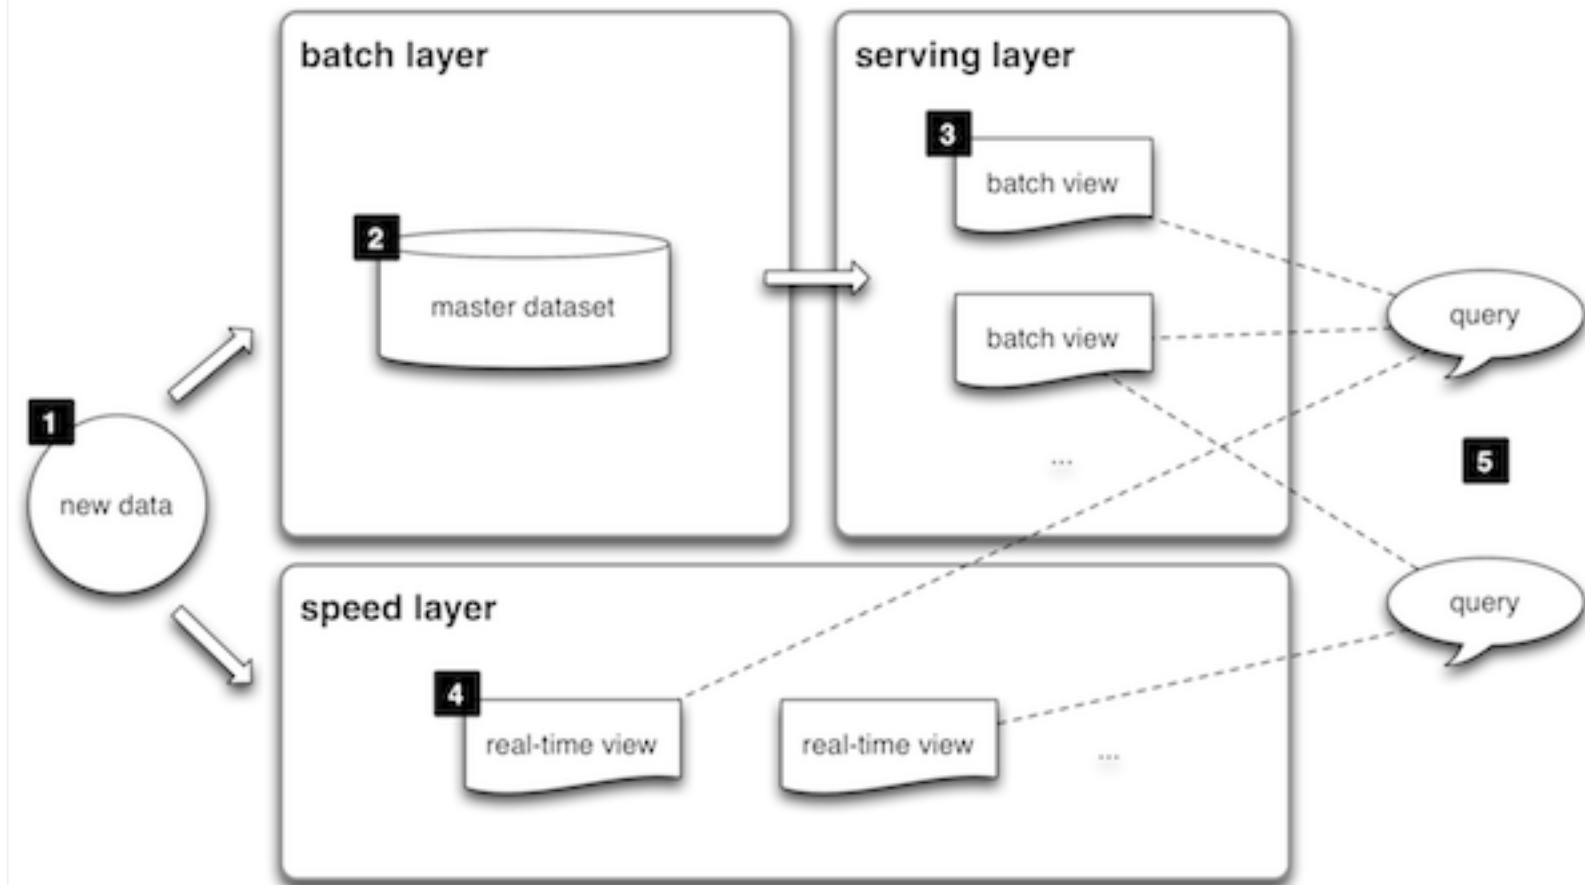
Define: Streaming

- A type of data processing engine that is designed with infinite data sets in mind
- Other common terms for streaming: unbounded data, low latency, approximate, and/or speculative results

Hurdles

Lambda architecture

Typical Lambda architecture



Latency

new tools...

Highly Scalable Blog

Articles on Big Data, NoSQL, and Highly Scalable Software Engineering

Probabilistic Data Structures for Web Analytics and Data Mining

Posted on May 1, 2012

<https://highlyscalable.wordpress.com/2012/05/01/probabilistic-structures-web-analytics-data-mining/>

We need new tools

```
SELECT avg(sessionTime)
FROM Table
WHERE city='San Francisco'
WITHIN 2 SECONDS
```

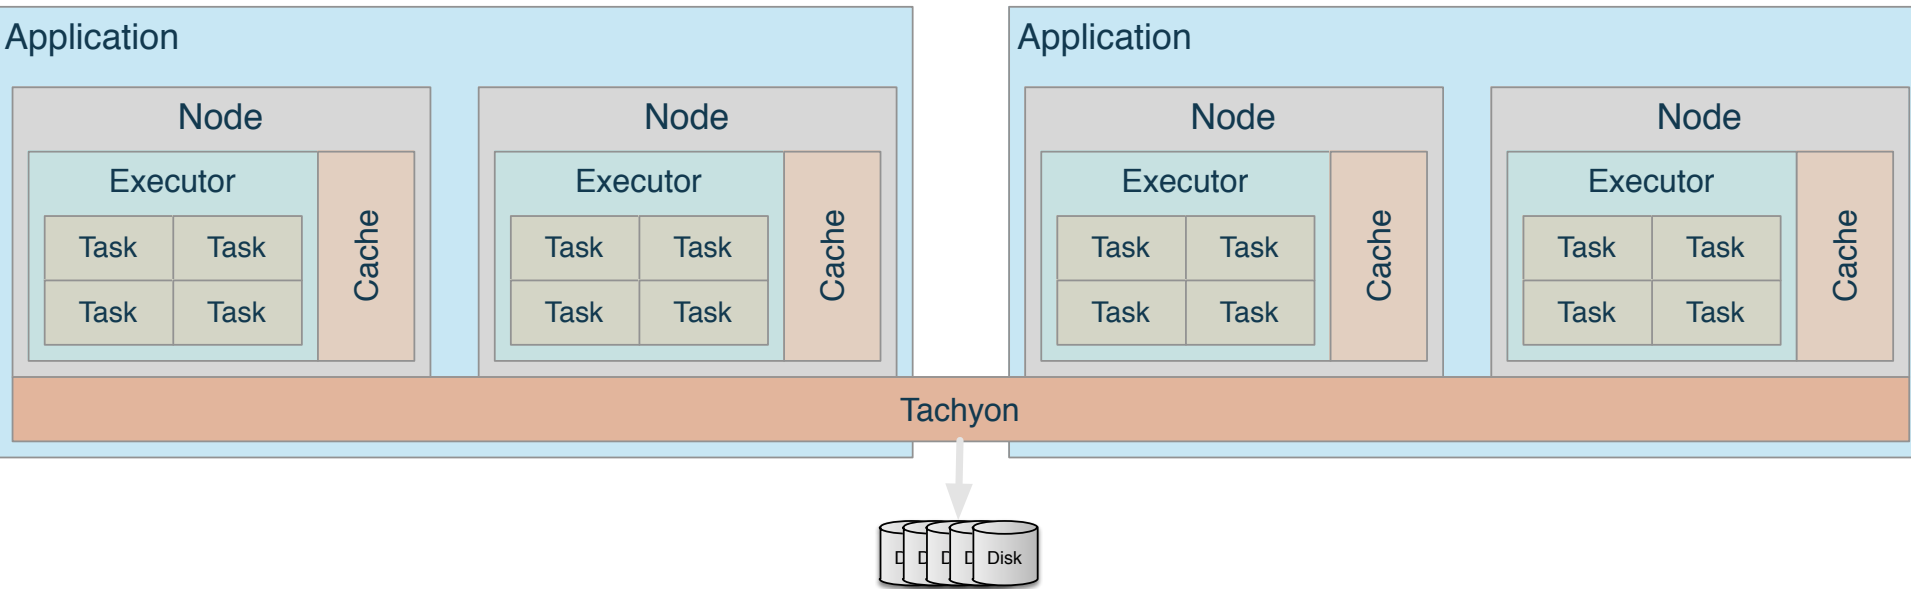
Queries with Time Bounds

```
SELECT avg(sessionTime)
FROM Table
WHERE city='San Francisco'
ERROR 0.1 CONFIDENCE 95.0%
```

Queries with Error Bounds



new tools...



Online ML



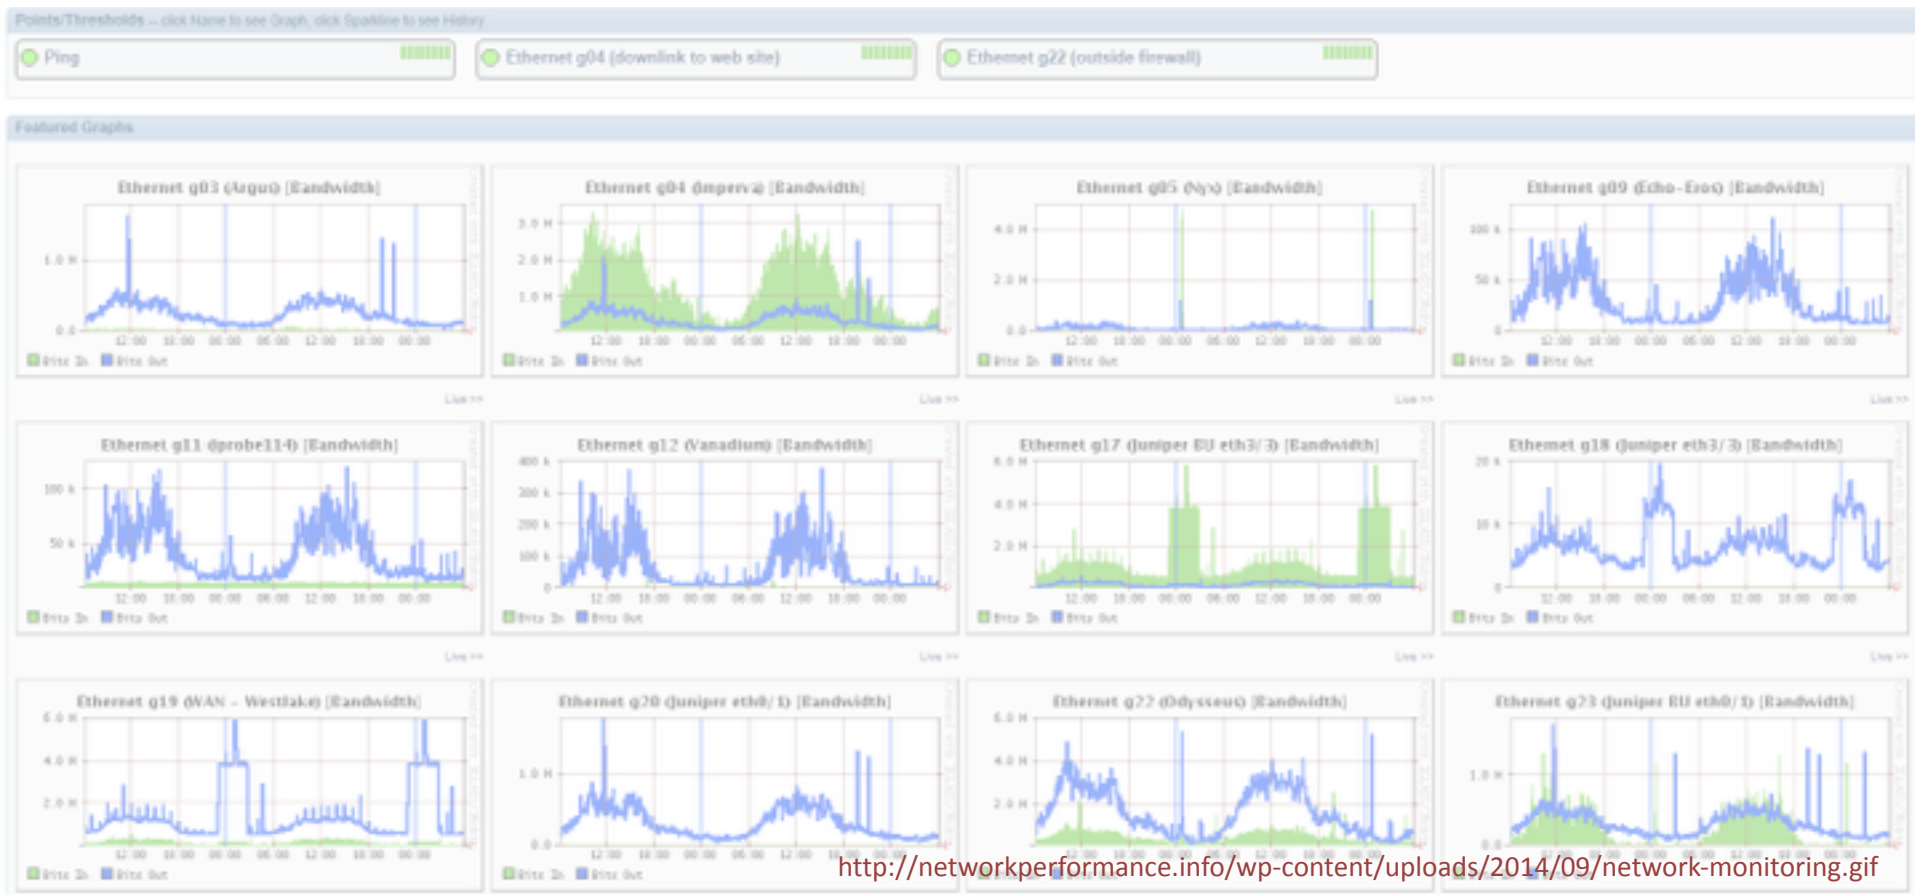
Computation model

```
public static class LineIndexMapper
    extends MapReduceBase
    implements Mapper<LongWritable, Text,
                    Text, Text> {
    private final static Text word =
        new Text();
    private final static Text location =
        new Text();

    public void map(
        LongWritable key, Text val,
        OutputCollector<Text, Text> output,
        Reporter reporter) throws IOException {

        FileSplit fileSplit =
            (FileSplit)reporter.getInputSplit();
```

Monitoring



Wish list

Streaming platform feature list

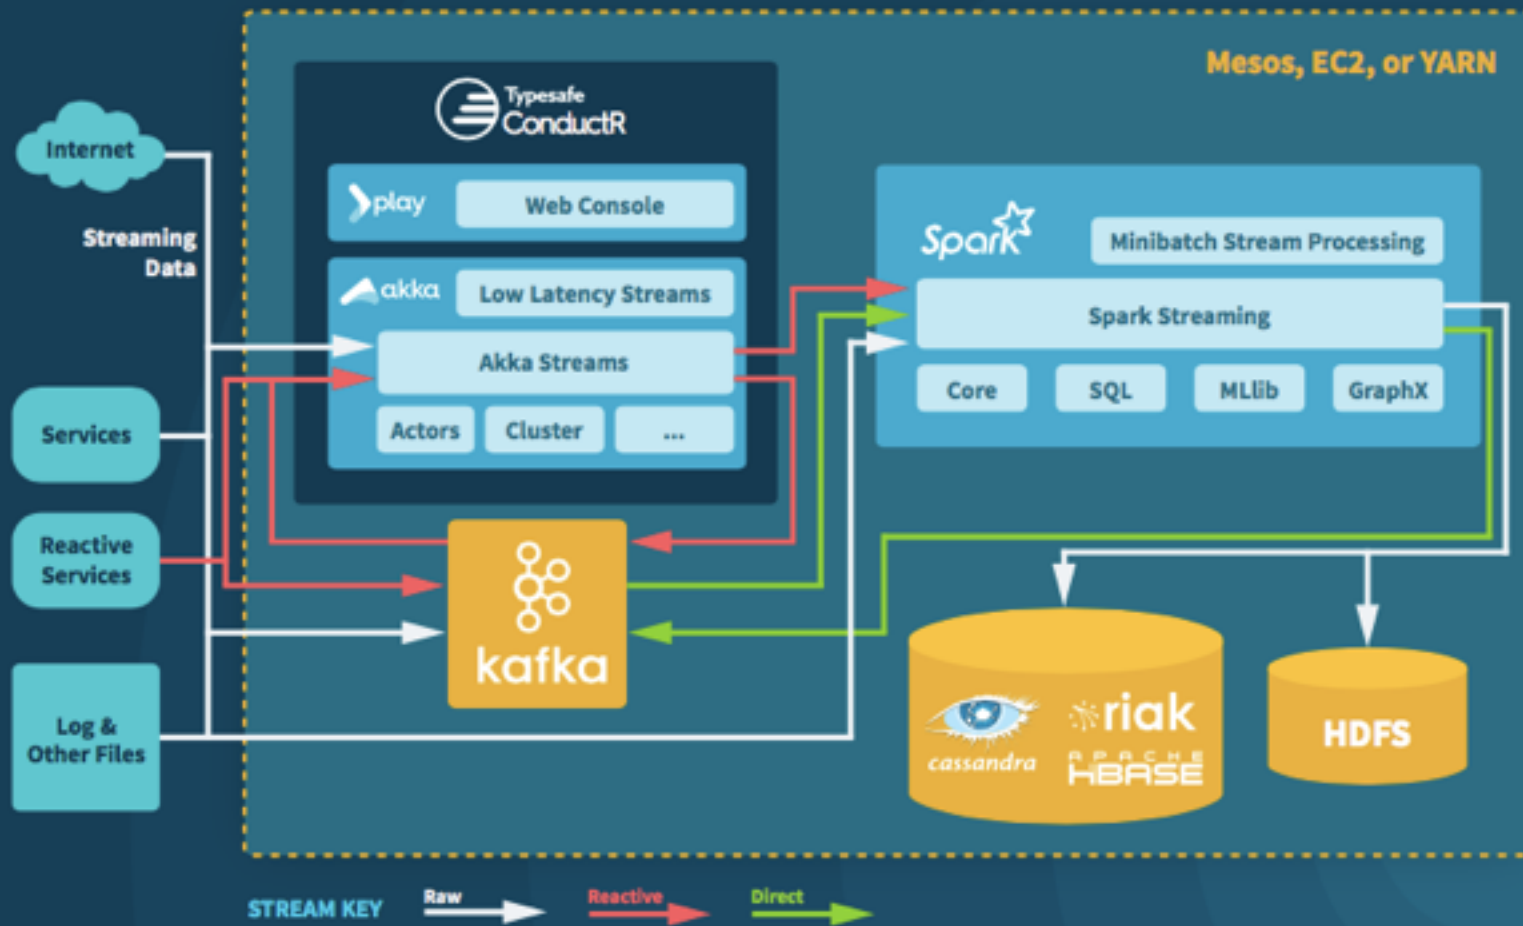
- Should be programmer friendly (Scala and FP)
- Should be fault tolerant
- Should support high throughput/low latency
- Should integrate with batch systems (Hadoop)
- Should run machine learning algorithm
- Should not overwhelm the consumer
- Should provide integration points for other streaming systems

And that streaming platform is...

not there yet

but...something is emerging

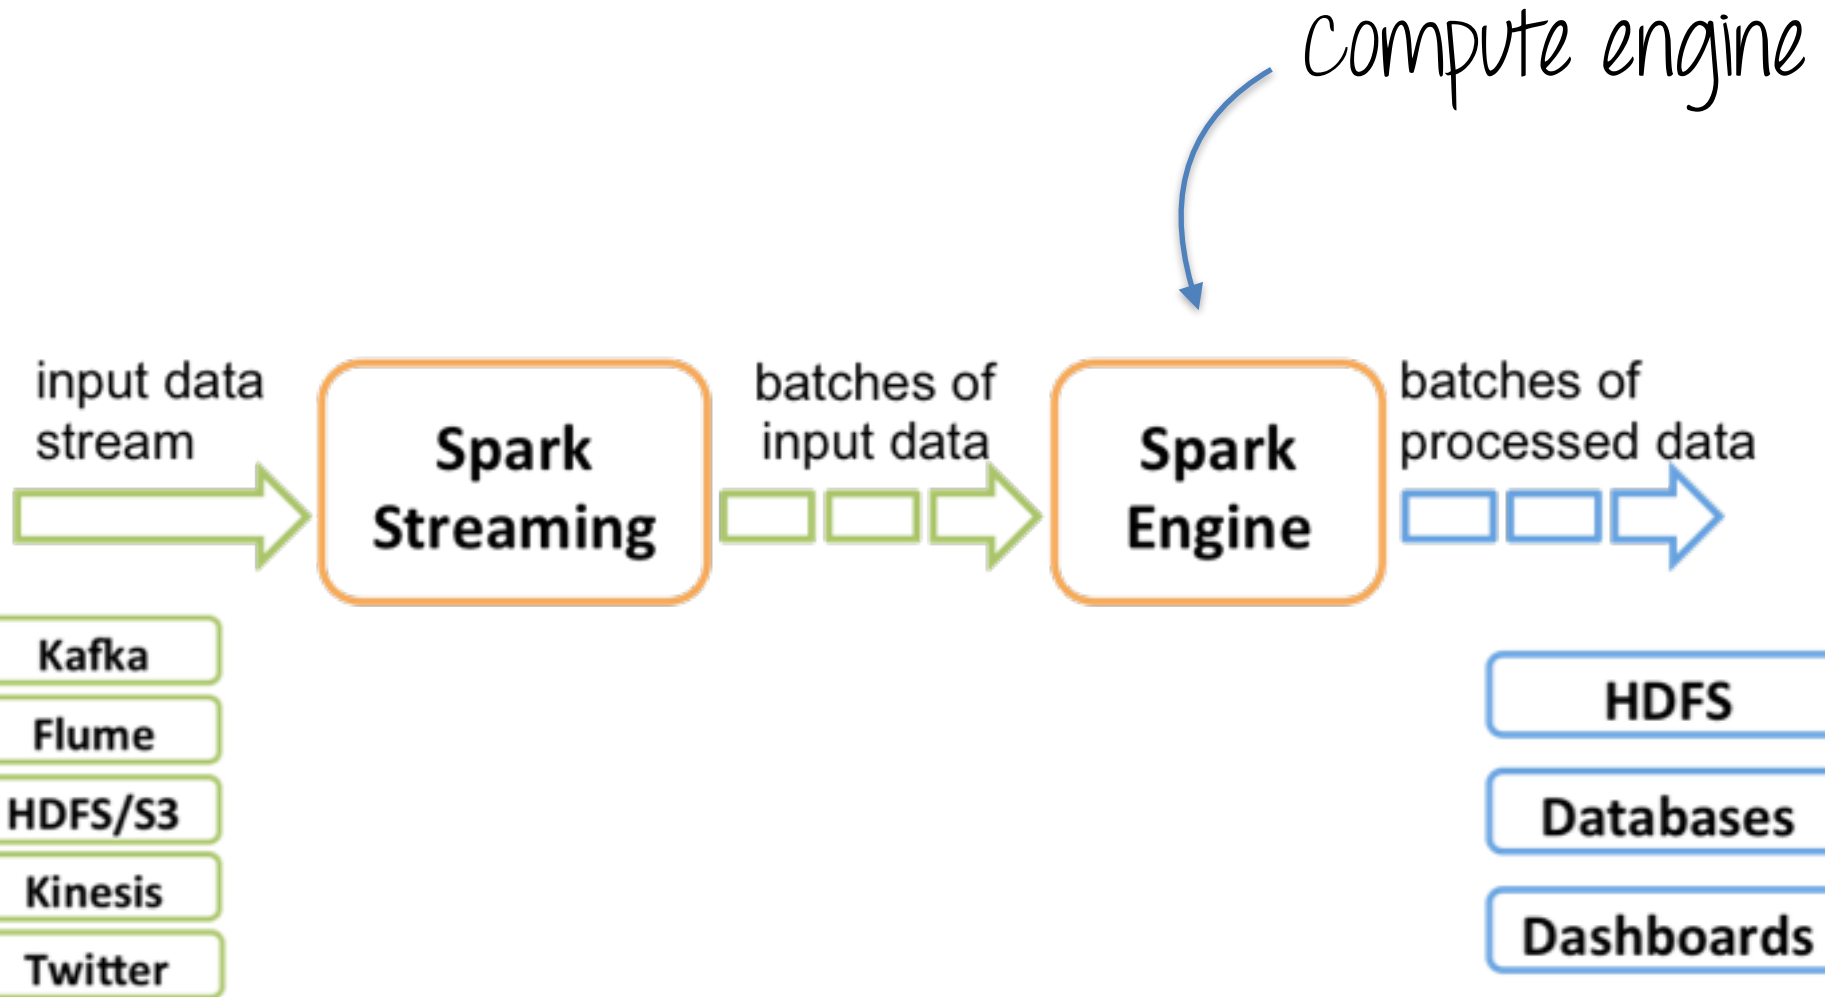
Fast data platform





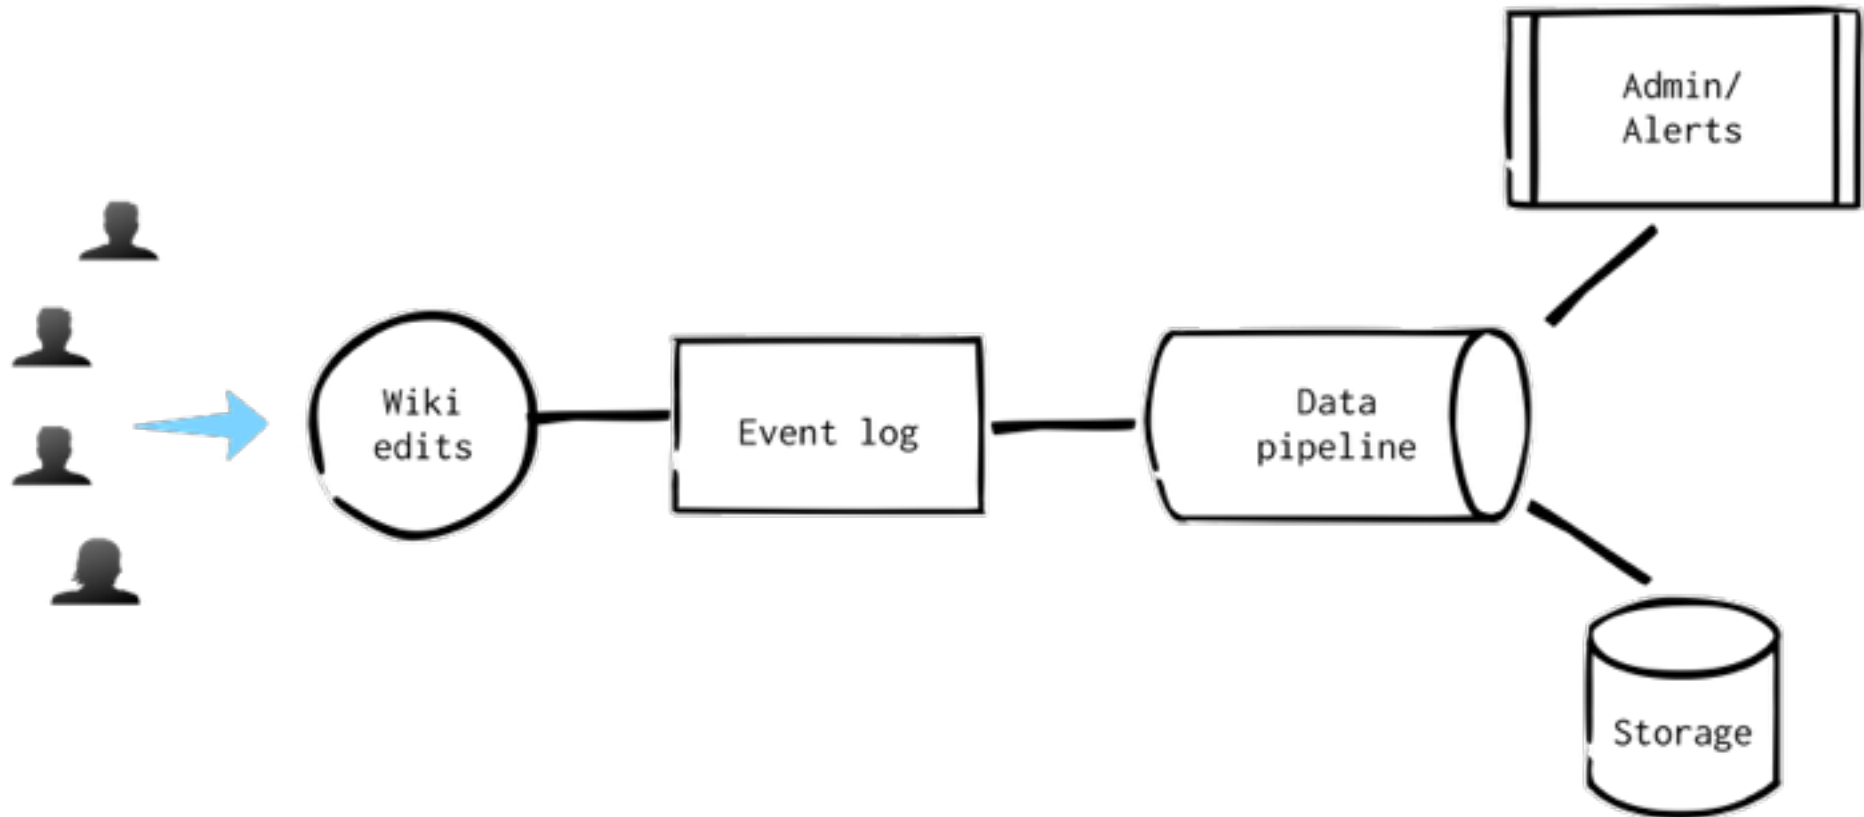
Captures time slices of events

Data pipeline

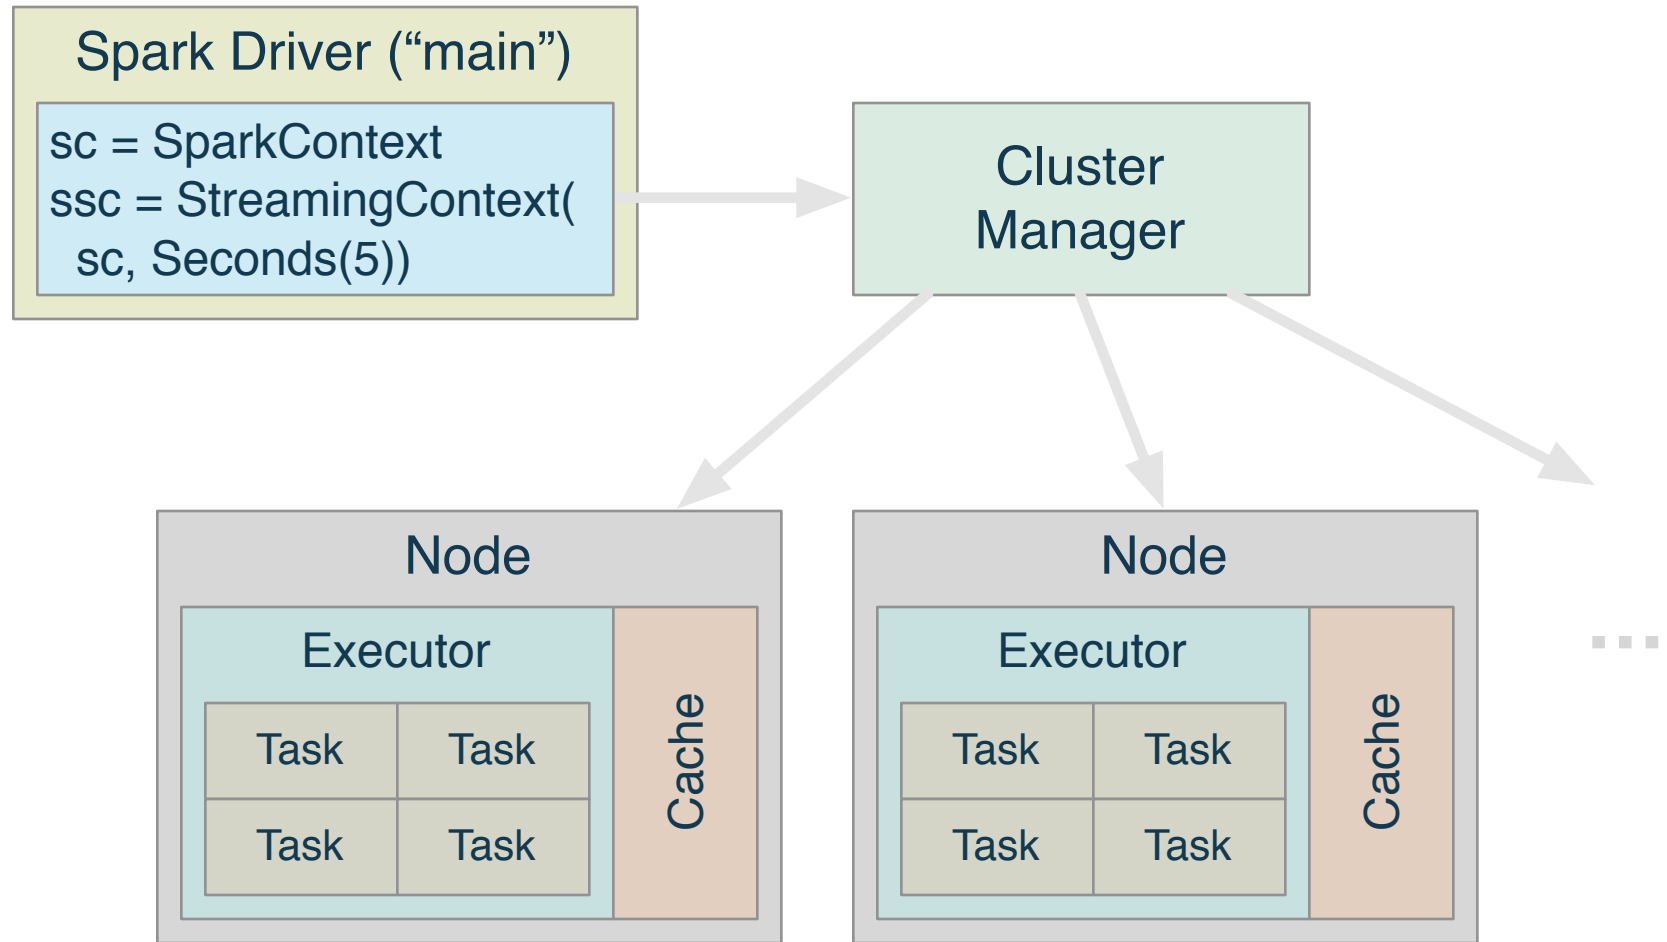


Show me the code

Predict breaking news



Under the hood




Reactive Streams

Definition


A standard for asynchronous stream processing with non-blocking back pressure

Supply and Demand




Dynamic Push-Pull

- “push” behavior when consumer is faster
- “pull” behavior when producer is faster




Explicit Demand: Tailored Flow Control



splitting the data means merging the demand

Explicit Demand: Tailored Flow Control

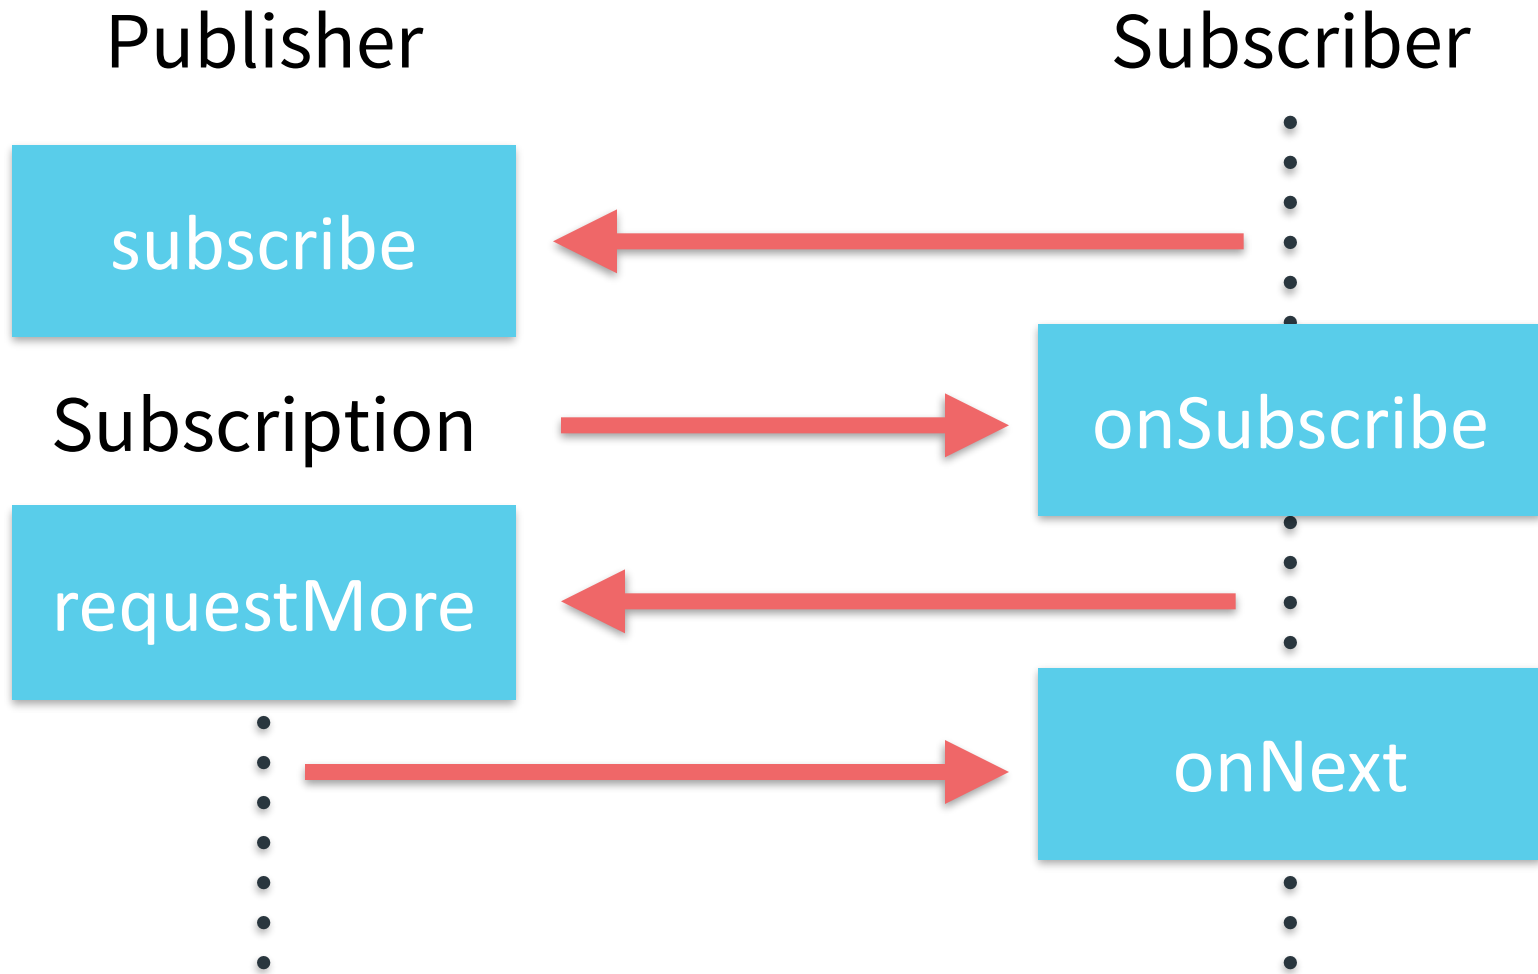


merging the data means splitting the demand

The Meat

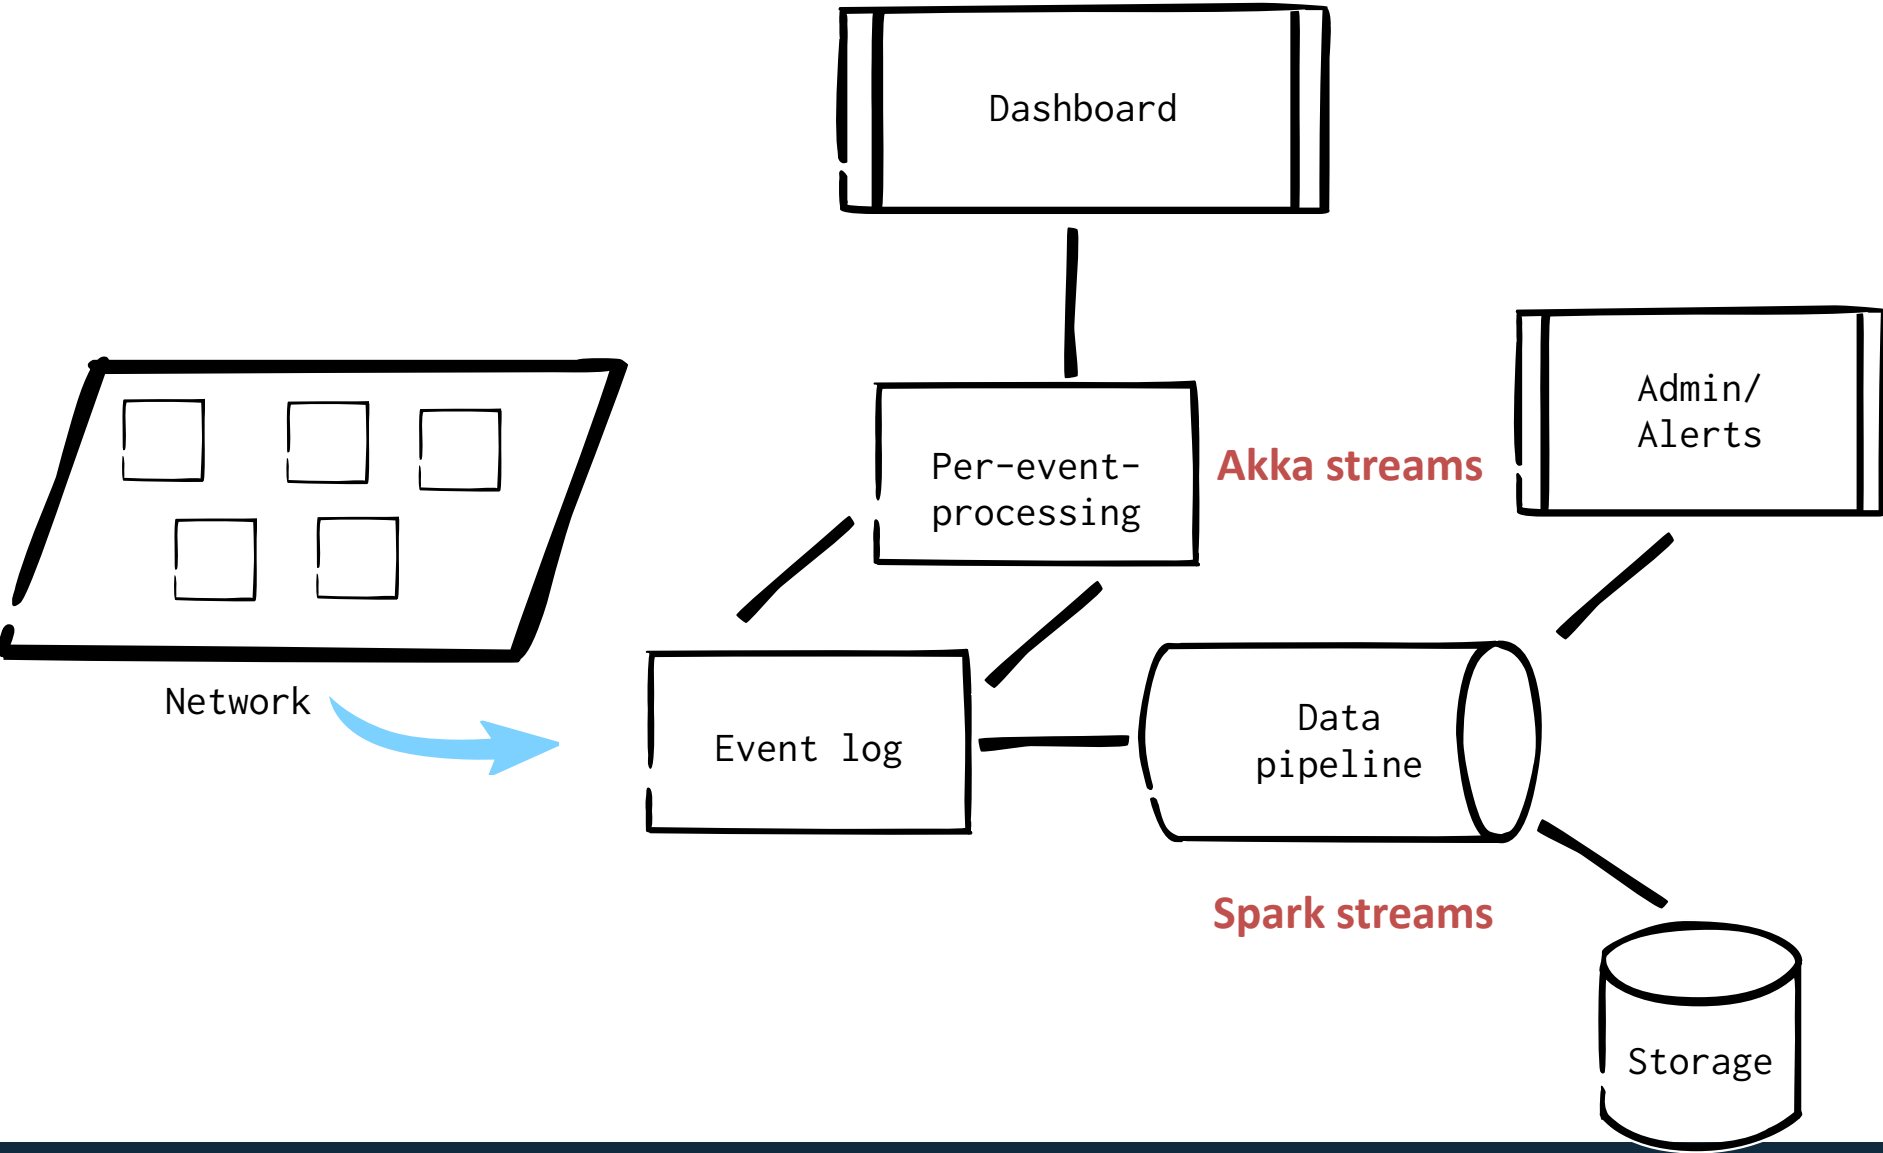
```
trait Publisher[T] {  
  def subscribe(sub: Subscriber[T]): Unit  
}  
  
trait Subscription {  
  def requestMore(n: Int): Unit  
  def cancel(): Unit  
}  
  
trait Subscriber[T] {  
  def onSubscribe(s: Subscription): Unit  
  def onNext(elem: T): Unit  
  def onError(thr: Throwable): Unit  
  def onComplete(): Unit  
}
```

How does it Connect?



Network intrusion

Detecting network intrusion



```
val sc = new SparkContext("local[3]", "Intro")  
val ssc = new StreamingContext(sc, Seconds(1))  
ssc.checkpoint("data/checkpoint")
```

```
val kmeans = new StreamingKMeans()  
    .setK(10)  
    .setRandomCenters(3, 100.0)  
    .setHalfLife(5, "batches")
```

```
val receiver = SparkEnv.get.actorSystem.actorOf(trainingDataReceiver,
  "training-data-source")
val trainingStream =
  ssc.actorStream[Vector](TrainingDataStream.props(receiver, kmeans),
    "training-stream")

val rawData: DStream[(String, Vector)] = KafkaUtils
  .createStream(ssc = ssc, kafkaParams, topics,
    StorageLevel.MEMORY_ONLY)
  .map(_._2).map(line => (line, labelAndVector(line)._2))
```

```
kmeans.trainOn(trainingStream.map(_._2))  
kmeans.predictOnValues(networkStream).print()  
  
ssc.start()  
ssc.awaitTermination()
```

Q & Option[A]

@nraychaudhuri



©Typesafe 2014 – All Rights Reserved